


[illegible]

Department of Computer Science and Electronics
Faculty of Mathematics and Natural Sciences
Universitas Gadjah Mada, Yogyakarta, Indonesia
Email: wahyo@ugm.ac.id



Graphs

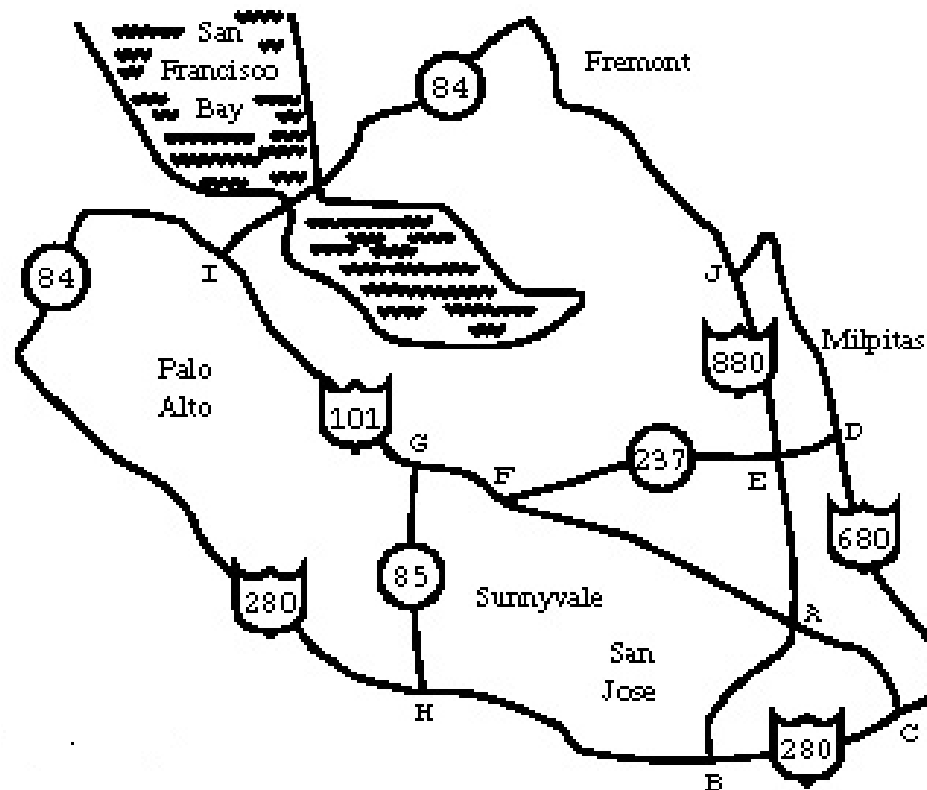
- Graphs are a data structure which represent relationships between entities
 - *Vertices* represent entities
 - *Edges* represent some kind of relationship





Example


- The graph on the previous page could be used to model San Jose freeway connections:





Adjacency


- Two vertices are *adjacent* to one another if they are connected by a single edge
- For example:
 - I and G are adjacent
 - A and C are adjacent
 - I and F are not adjacent
- Two adjacent nodes are considered *neighbors*





Path

- A *path* is a sequence of edges




- Paths in this graph include:
 - BAEJ, CAFG, HGFEJDCAB, HIJDCBAFE, etc.



Connected Graphs


- A graph is connected if there is at least one path from every vertex to every other vertex





Unconnected Graph

- An *unconnected graph* consists of several *connected components*:




- Connected components of this graph are:
 - AB, and CD
- We'll be working with connected graphs



Directed Graphs


- A graph where edges have directions
 - Usually designated by an arrow





Weighted Graphs

- A graph where edges have weights, which quantifies the relationship
 - For example, you may assign path distances between cities
 - Or airline costs
- These graphs can be directed or undirected






Vertices: Java Implementation

- We can represent a vertex as a Java class with:
 - Character data
 - A boolean data member to check if it has been visited
- Now we need to specify edges.
 - But in a graph, we don't know how many there will be!
- We can do this using either an adjacency *matrix* or an adjacency *list*. We'll look at both.



Adjacency Matrix

- An adjacency matrix for a graph with n nodes, is size $n \times n$
 - Position (i, j) contains a 1 if there is an edge connecting node i with node j
 - Zero otherwise
- For example, here is a graph and its adjacency matrix:




	A	B	C	D
A	0	1	1	1
B	1	0	0	1
C	1	0	0	0
D	1	1	0	0



Redundant?

- This may seem a bit redundant:




	A	B	C	D
A	0	1	1	1
B	1	0	0	1
C	1	0	0	0
D	1	1	0	0

- Why store two pieces of information for the same edge?
 - i.e., (A, B) and (B, A)
- Unfortunately, there's no easy way around it
 - Because edges have no direction
 - No concept of 'parents' and 'children'



Adjacency List

- An array of linked lists
 - Index by vertex, and obtain a linked list of neighbors
- Here is the same graph, with its adjacency list:



Vertex	List Containing Adjacent Vertices
A	B—>C—>D
B	A—>D
C	A
D	A—>B




Application: Searches

- A fundamental operation for a graph is:
 - Starting from a particular vertex
 - Find all other vertices which can be reached by following paths
- Example application
 - How many towns in the US can be reached by train from Tampa?
- Two approaches
 - Depth first search (DFS)
 - Breadth first search (BFS)



Depth First Search (DFS)

- Idea
 - Pick a starting point
 - Follow a path to unvisited vertices, as long as you can until you hit a dead end
 - When you hit a dead end, go back to a previous spot and hit unvisited vertices
 - Stop when every path is a dead end






Depth First Search (DFS)

- Algorithm
 - Pick a vertex (call it A) as your starting point
 - Visit this vertex, and:
 - Push it onto a stack of visited vertices
 - Mark it as visited (so we don't visit it again)
 - Visit any neighbor of A that hasn't yet been visited
 - Repeat the process
 - When there are no more unvisited neighbors
 - Pop the vertex off the stack
 - Finished when the stack is empty
- Note: We get as far away from the starting point until we reach a dead end, then pop (can be applied to mazes)



Example

- Start from A, and execute depth first search on this graph, showing the contents of the stack at each step
 - Every step, we'll either have a visit or pop





Depth First Search: Complexity

- Let $|V|$ be the number of vertices in a graph
- And let $|E|$ be the number of edges
- In the worst case, we visit every vertex and every edge:
 - $O(|V| + |E|)$ time


At first glance, this doesn't look too bad

- But remember a graph can have lots of edges!
- Worst case, every vertex is connected to every other:
 - $(n-1) + (n-2) + \dots + 1 = O(n^2)$
- So it can be expensive if the graph has many edges

Breadth First Search (BFS)



- Same application as DFS; we want to find all vertices which we can get to from a starting point, call it A
- However this time, instead of going as far as possible until we find a dead end, like DFS
 - We visit all the closest vertices first
 - Then once all the closest vertices are visited, branch further out






Breadth First Search (BFS)

- We're going to use a queue instead of a stack!
- Algorithm
 - Start at a vertex, call it current
 - If there is an unvisited neighbor, mark it, and insert it into the queue
 - If there is not:
 - If the queue is empty, we are done, otherwise:
 - Remove a vertex from the queue and set current to that vertex, and repeat the process



Example

- Start from A, and execute breadth first search on this graph, showing the contents of the queue at each step
 - Every step, we'll either have a visit or a removal






Breadth First Search: Complexity

- Let $|V|$ be the number of vertices in a graph
- And let $|E|$ be the number of edges
- In the worst case, we visit every vertex and every edge:
 - $O(|V| + |E|)$ time
 - Same as DFS
- Again, if the graph has lots of edges, you approach quadratic run time which is the worst case


Minimum Spanning Trees (MSTs)



- On that note of large numbers of edges slowing down our precious search algorithms:
 - Let's look at MSTs, which can help ameliorate this problem
- It would be nice to take a graph and reduce the number of edges to the minimum number required to span all vertices:



a) Extra Edges




b) Minimum Number of Edges

What's the number of edges now?




We've done it already...

- Actually, if you execute DFS you've already computed the MST!



a) Extra Edges




b) Minimum Number of Edges

- Think about it: you follow a path for as long as you can, then backtrack (visit every vertex at most once)
 - You just have to save edges as you go



Directed Graphs

- A directed graph is a graph where the edges have direction, signified by arrows:




- This will simplify the adjacency matrix a bit...



Adjacency Matrix

- The adjacency matrix for this graph does not contain redundant entries
 - Because now each edge has a source and a sink
 - So entry (i, j) is only set to 1 if there is an edge going from i to j
 - 0 otherwise




	A	B	C
A	0	1	0
B	0	0	1
C	0	0	0

Topological Sort



- Only works with DAGs (Directed Acyclic Graphs)
 - That is if the graph has a cycle, this will not work




- Idea: Sort all vertices such that if a vertex's successor always appears after it in a list (application: course prerequisites)




Topological Sorting

- Start from vertex that have in-degree 0






Topological Sorting






Topological Sorting






Topological Sorting






Topological Sorting






Topological Sorting






Topological Sorting






Topological Sorting






Topological Sorting





Topological Sorting





UNIVERSITAS GADJAH MADA

THANK YOU

